

The PCA

Helwig Hauser *et al.*,
UiB Dept. of Informatics



Looking Back & Forth



Last time:

- from changing bases to the SVD
- SVD facts and interpreting the SVD
- data matrices and the SVD
- rank- k approximation with the SVD

Today:

- variance, covariance, and correlation
- eigenanalysis
- PCA (principal component analysis)


Principal Component Analysis (PCA):

- represents the **data** (also) **in an alternative frame** such that *most data variance is aligned with the new "x-axis"*
- PCA is **based on an eigenanalysis of the covariance matrix** and the corresponding *eigenvalues indicate, how much variance is explained by each axis*
- PCA can be used as a *basis for* **dimension reduction**: if only those k axes are used that correspond to the largest variance in the data, this leads to a k -dimensional approximation
- PCA can be useful, *when dealing with noisy data*: much of the signal can be concentrated into the first components, possibly raising the SNR (after dimension reduction)

Variance

In 1D:

- 1st and 2nd order *moments* in descriptive statistics:
 - the *mean* μ estimates the *center* of the data:
 - the *standard deviation* σ estimates the *variation / dispersion* of the data (68–95–99.7–rule):
 - the *variance* σ^2 is the *square of the standard deviation*:



$$\sigma = \sqrt{\frac{1}{n} \sum_i (x_i - \mu)^2}$$

$$\sigma^2 = \frac{1}{n} \sum_i (x_i - \mu)^2$$

Rewriting variance:

- shift the data by $-\mu$ to center it
- computing the variance as $\mathbf{x}^T \mathbf{x} / (n-1)$ (col. \mathbf{x})
- $\mathbf{x} = [65; 58; 74; 54; 72; 62; 72]$

!

```
mu = mean(x)
```

```
x_c = x - mu
```

```
x_size = size(x); n = x_size(1)
```

```
sigma_2 = (x_c.' * x_c) / (n-1)
```

note:
for an unbiased
estimator,
divide by $n-1$

Variance and Covariance

In $n\mathbf{D}$:

- the center is estimated per dimension (one may also rescale!)
- the variation is estimated (after centering):
 - within each dimension \mathbf{x} : $\sigma^2_{\mathbf{x}} = \mathbf{x}^T \mathbf{x} / (n-1)$
 - between dims. \mathbf{x} & \mathbf{y} (covariance): $\sigma^2_{\mathbf{xy}} = \mathbf{x}^T \mathbf{y} / (n-1)$

All variances / covariances in $n\mathbf{D}$:

- variances and covariances show up in the covariance matrix:
 - given data matrix \mathbf{D} (items in cols.): $\Sigma^2_{\mathbf{D}} = \mathbf{D} \mathbf{D}^T / (n-1)$
 - transposed matrix \mathbf{D} (items in rows): $\Sigma^2_{\mathbf{D}} = \mathbf{D}^T \mathbf{D} / (n-1)$

Eigenanalysis





Assume a square matrix \mathbf{A} in $\mathbb{R}^{n \times n}$:

- then an *eigenvalue problem* is written as $\mathbf{A} \mathbf{x} = \lambda \mathbf{x}$ with λ being a scalar value, i.e., an *eigenvalue* of \mathbf{A}
- to solve an eigenvalue problem for both \mathbf{x} and λ means to find—in addition to all **eigenvalues** λ_i —all corresponding vectors \mathbf{x}_i , i.e., the **eigenvectors** \mathbf{x}_i , for which the left-multiplication with \mathbf{A} is the same as a scalar stretching/shrinking of \mathbf{x} (by λ) (eigenvectors \mathbf{x} do not change direction under \mathbf{A} , only length, if at all)

[ex. from Wikipedia]

the blue vector is an eigenvector
of the depicted shear operation
(but not so the red vector)



Eigenvalues



Given a square, real-valued matrix \mathbf{A} , then:

- the eigenvalues can be positive, zero, or negative
- the eigenvalues can be complex (they appear in pairs for $\mathbf{A} \in \mathbb{R}^{n \times n}$)
- the eigenvalues can appear multiple times (algebraic multiplicity μ)
- all eigenvalues add up to the trace of \mathbf{A}
- all eigenvalues multiply to the determinant of \mathbf{A}
- \mathbf{A} is invertible, *if and only if* all eigenvalues of \mathbf{A} are non-zero
- \mathbf{A} is invertible \Rightarrow the eigenvalues of \mathbf{A}^{-1} are $1/\lambda_i$
- if \mathbf{A} is unitary, then the norm of all eigenvalues of \mathbf{A} is 1
- if $\mathbf{A} = \mathbf{A}^*$ (\mathbf{A} is self-adjoint), then all eigenvalues are real (in particular true for symmetric, real-valued matrices!)
- the eigenvalues of \mathbf{A}^k are λ_i^k , if k a positive integer
- if \mathbf{A} is positive-/negative (semi-)definite, then all eigenvalues of \mathbf{A} are positive/negative (incl. 0)

Eigenvectors



Each eigenvalue λ_i (algebraic multiplicity $\mu_i \geq 1$)

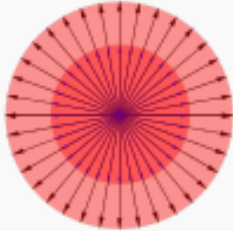
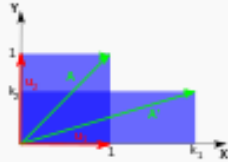
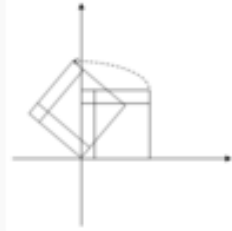
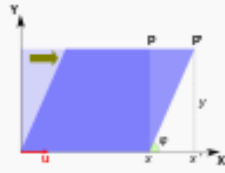
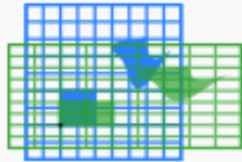
- corresponds to a subspace, i.e., the *eigenspace spanned by the corresponding eigenvector(s)*
- if **A** is symmetric,
then the eigenvectors are orthogonal to each other
- any multiple of an eigenvector is also an eigenvector
- the eigenvectors can be complex (pairs of them)
- eigenspaces are γ_i -dimensional ($\gamma_i \geq 1$) – geometric multiplicity (the usual case is 1D)

in MATLAB: „ $[V, L]=\text{eig}(A)$ “ gives
the (normalized) eigenvectors (as columns of **V**) and
the eigenvalues (as major diagonal of **L**)

Eigenvalues and Eigenvectors in 2D



[from Wikipedia]

	scaling	unequal scaling	rotation	horizontal shear	hyperbolic rotation
illustration					
matrix	$\begin{bmatrix} k & 0 \\ 0 & k \end{bmatrix}$	$\begin{bmatrix} k_1 & 0 \\ 0 & k_2 \end{bmatrix}$	$\begin{bmatrix} c & -s \\ s & c \end{bmatrix}$ $c = \cos \theta$ $s = \sin \theta$	$\begin{bmatrix} 1 & k \\ 0 & 1 \end{bmatrix}$	$\begin{bmatrix} c & s \\ s & c \end{bmatrix}$ $c = \cosh \varphi$ $s = \sinh \varphi$
characteristic polynomial	$(\lambda - k)^2$	$(\lambda - k_1)(\lambda - k_2)$	$\lambda^2 - 2c\lambda + 1$	$(\lambda - 1)^2$	$\lambda^2 - 2c\lambda + 1$
eigenvalues λ_i	$\lambda_1 = \lambda_2 = k$	$\lambda_1 = k_1$ $\lambda_2 = k_2$	$\lambda_1 = e^{i\theta} = c + si$ $\lambda_2 = e^{-i\theta} = c - si$	$\lambda_1 = \lambda_2 = 1$	$\lambda_1 = e^{\varphi}$ $\lambda_2 = e^{-\varphi}$
algebraic multipl. $\mu_i = \mu(\lambda_i)$	$\mu_1 = 2$	$\mu_1 = 1$ $\mu_2 = 1$	$\mu_1 = 1$ $\mu_2 = 1$	$\mu_1 = 2$	$\mu_1 = 1$ $\mu_2 = 1$
geometric multipl. $\gamma_i = \gamma(\lambda_i)$	$\gamma_1 = 2$	$\gamma_1 = 1$ $\gamma_2 = 1$	$\gamma_1 = 1$ $\gamma_2 = 1$	$\gamma_1 = 1$	$\gamma_1 = 1$ $\gamma_2 = 1$
eigenvectors	All non-zero vectors	$u_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$ $u_2 = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$	$u_1 = \begin{bmatrix} 1 \\ -i \end{bmatrix}$ $u_2 = \begin{bmatrix} 1 \\ +i \end{bmatrix}$	$u_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$	$u_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$ $u_2 = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$

Writing it all-in-one: a Decomposition!




Assuming n solutions \mathbf{x}_i & λ_i to $\mathbf{A} \mathbf{x} = \lambda \mathbf{x}$ (multiplicity OK)

$$\begin{array}{l} \text{– then} \quad \left. \begin{array}{l} \mathbf{A} \mathbf{x}_1 = \lambda_1 \mathbf{x}_1 \\ \mathbf{A} \mathbf{x}_2 = \lambda_2 \mathbf{x}_2 \\ \vdots \\ \mathbf{A} \mathbf{x}_n = \lambda_n \mathbf{x}_n \end{array} \right\} \Rightarrow \mathbf{A} \mathbf{V} = \mathbf{V} \mathbf{\Lambda} \end{array}$$

with $\mathbf{V} = (\mathbf{x}_1 \ \mathbf{x}_2 \ \dots \ \mathbf{x}_n)$ and

$$\mathbf{\Lambda} = \begin{pmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \lambda_n \end{pmatrix}$$


$$\Rightarrow \mathbf{A} = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^{-1}$$

A better Frame



Eigenanalysis...

- ... leads to the eigenvalue/eigenvector decomposition $\mathbf{A} = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^{-1}$ with \mathbf{V} hosting all (normalized) eigenvectors and $\mathbf{\Lambda}$ having all eigenvalues along the major diagonal
- ... if \mathbf{A} is real-valued and symmetric, then \mathbf{V} is unitary, i.e., an orthonormal basis of \mathbb{R}^n (important for PCA)

This suggests...

- ... that we can consider our \mathbf{A} in terms of basis \mathbf{V} : $\mathbf{X} = \mathbf{V}^* \mathbf{A}$ (\mathbf{X} is then \mathbf{A} in orthonormal \mathbf{V} -coordinates)

This leads to:

[[see the MATLAB example 2d_C](#)]

- Given \mathbf{A} in orthonormal \mathbf{V} -coords. (by $\mathbf{X} = \mathbf{V}^* \mathbf{A}$), $\mathbf{A} = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^*$ leads to $\mathbf{X} = \mathbf{V}^* \mathbf{A} = \mathbf{V}^* \mathbf{V} \mathbf{\Lambda} \mathbf{V}^* = \mathbf{\Lambda} \mathbf{V}^*$
- in particular $\mathbf{X} \mathbf{v} = \mathbf{\Lambda} \mathbf{V}^* \mathbf{v} = \mathbf{\Lambda} [\mathbf{v}]_{\mathbf{V}}$ (dim.-wise scaling!)

Eigenanalysis – Interpretation


Assuming a mapping $\mathbf{A}: \mathbb{R}^2 \rightarrow \mathbb{R}^2$ (again),

– and also considering

four vectors $\begin{pmatrix} 1 \\ 1 \end{pmatrix}, \begin{pmatrix} 2 \\ 1 \end{pmatrix}, \begin{pmatrix} 1 \\ 2 \end{pmatrix}, \begin{pmatrix} 2 \\ 2 \end{pmatrix}$

('*' below after the mapping)

$$\mathbf{A} = \begin{pmatrix} -0.4 & 1.2 \\ 1.2 & -1.1 \end{pmatrix}$$



Then the eigenanalysis of \mathbf{A} gives:

in MATLAB: „ $[\mathbf{V}, \mathbf{L}] = \text{eig}(\mathbf{A})$ “

– eigenvalues -2 and $\frac{1}{2}$

– and eigenvectors $\begin{pmatrix} -0.6 \\ 0.8 \end{pmatrix}^{\mathbf{v}_1}$ & $\begin{pmatrix} -0.8 \\ -0.6 \end{pmatrix}^{\mathbf{v}_2}$

[see the MATLAB example 2d_D]

SVD–Eigenanalysis Comparison



Both are decomposition/diagonalization approaches,

- i.e., both lead to a factorization of **A** into **P M Q** with **M** (singular values or eigenvalues) being diagonal and **P** & **Q** representing two change of basis steps

But


- the SVD exists always (+), even for rectangular matrices (+), while the eigendecomposition not necessarily exists (–)
- the SVD example: *nilpotent* matrix $\begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix}$ results in a real-valued solution (+), while the eigenvalues/eigenvectors may be complex-valued (–)
- the SVD uses two bases **U** & **V** (–), which always are unitary/orthonormal (+), the eigendecomp. uses one (+), **V**, that may not be orthogonal (–)

PCA – Introduction



Some data

- may be “intrinsically” r -dimensional, but represented d -dimensional with $d > r$
- example: if all data $\mathbf{d}_i \in \mathbb{R}^3$ are of form $\mathbf{d}_i = s_i \mathbf{v}$, $\mathbf{v} \in \mathbb{R}^3$, then their representation is 3D, while their intrinsic dimensionality is 1



Let's look at some people's data:

- we consider Height vs. Shoesize now (they look correlated, right?!)
- the covariance matrix looks like

$$\Sigma_D^2 = \begin{pmatrix} 388.7857 & \underline{82.9229} \\ \underline{82.9229} & 20.3942 \end{pmatrix}$$

with major off-diagonal values!

[see the MATLAB example 2d_E]

PCA – Diagonalizing Σ^2

Next step (after normalization/standardization):

- transform such that the covariance matrix becomes diagonal!
(since Σ^2 is real-valued and symmetric, we use the eigendecomp.)


This leads to the eigenvalues & eigenvectors of Σ^2 :

- eigenvalues ≈ 1.93 and ≈ 0.07
with the corresponding eigenvectors

$$\mathbf{V} = \begin{pmatrix} -0.7071 & 0.7071 \\ -0.7071 & -0.7071 \end{pmatrix}$$

called *principal components* (PC)

- the vast difference in eigenvalues tells that this dataset is mostly 1D!



[see the MATLAB example 2d_E]

PCA – Look at the Scores




Next step:

- we can now look at the data in terms of \mathbf{V} , i.e., in the PC of \mathbf{D} (the $[\mathbf{D}]_{\mathbf{V}}$ are called the *scores* in PCA)

The scores are then uncorrelated

- their covariance matrix is diagonal
- the biggest data variation happens along \mathbf{v}_1 , followed by \mathbf{v}_2 , etc.




PCA – Low-rank Approximation



Similar to the SVD,

- we can compute a lower- $(k-)$ dimensional approximation of the (normalized) data \mathbf{D} by “only” considering the largest eigenvalues



Given scores $[\mathbf{D}]_{\mathbf{v}_i}$

- we keep only the data rows corresponding to the largest k eigenvalues
- assuming that we reordered \mathbf{V} & $\mathbf{\Lambda}$ so that the eigenvalues in $\mathbf{\Lambda}$ became sorted in decreasing size
- we then compute \mathbf{D}'_k by
$$\mathbf{D}'_k = (\mathbf{v}_1 \mathbf{v}_2 \dots \mathbf{v}_k) (\mathbf{d}_{\mathbf{v}_1}^T \mathbf{d}_{\mathbf{v}_2}^T \dots \mathbf{d}_{\mathbf{v}_k}^T)^T$$

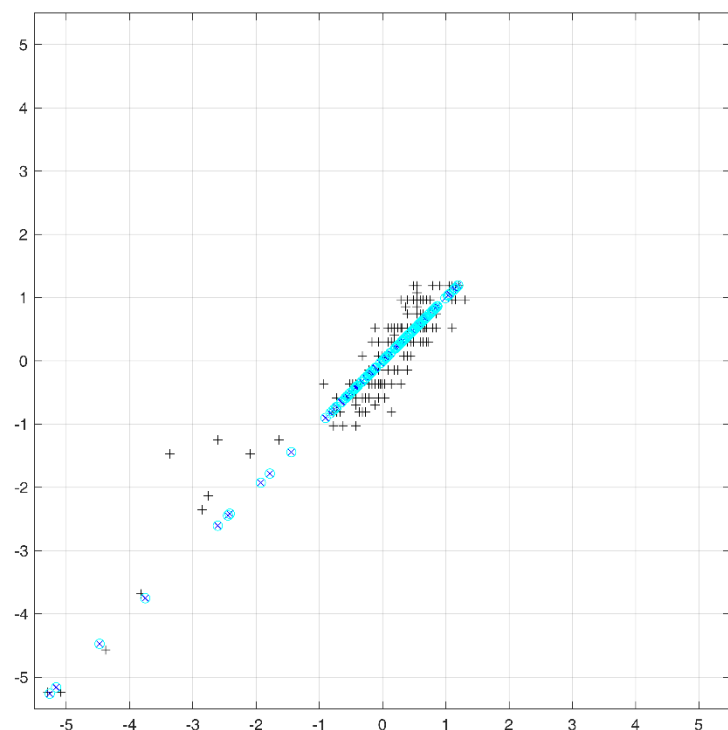
[see the MATLAB example 2d_E]

PCA vs. SVD



The lower-dimensional approximation with PCA overlaps with the rank- k approximation with SVD:

- the eigenvalues from PCA, λ_i , are equal to $\sigma_i^2/(n-1)$
- the left-singular vectors \mathbf{U} are the PC from PCA (up to sign change), if data rows are given (otherwise it's \mathbf{V})



Relation:

1. data \mathbf{D} , d dims. \times n items (in the cols.)
2. $\mathbf{\Sigma}^2 = \mathbf{D} \mathbf{D}^T / (n-1)$ // covariance-matrix
3. $\mathbf{\Sigma}^2 = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^*$ // eigendecomposition
4. $[\mathbf{D}]_v = \mathbf{V}^* \mathbf{D}$ // scores
5. $\mathbf{D} = \mathbf{U} \mathbf{S} \mathbf{W}^*$ // SVD
6. $\mathbf{\Sigma}^2 = \mathbf{U} \mathbf{S} \mathbf{W}^* (\mathbf{W} \mathbf{S}^* \mathbf{U}^*) / (n-1)$ // 2. & 5.
 $= \mathbf{U} (\mathbf{S} \mathbf{S}^* / (n-1)) \mathbf{U}^* = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^*$ // 3.

PCA – Investigate the Loadings (1)



With PCA

- we transform the data into a new Cartesian coordinate frame (the PCs are orthogonal to each other)
- so that the covariance between the scores is zero (we remove any redundancy from the data)

Challenging (as with the SVD):

- considering the scores $[\mathbf{D}]_{\mathbf{v}}$ of data \mathbf{D} , esp. after normalization, we usually lack a good interpretation of the new axes—they are linear combinations of the original data-axes, i.e., the “new data” $[\mathbf{D}]_{\mathbf{v}}$ is a mix (linear combination) of \mathbf{D}
- one way to interpret the principal components, is to look at their *loadings*:
to which degree play the original data-axes into them?

PCA – Investigate the Loadings (2)



The loadings

- of PC i are given by the components of eigenvector \mathbf{v}_i
- given data rows (items in cols.),
 - data item j is given by column j in \mathbf{D} : $(d_{...j})$
 - the scores of data item j are then $(d_{\mathbf{v}_{...j}}) = \mathbf{V}^* (d_{...j})$,
i.e., $d_{\mathbf{v}_{1j}} = \mathbf{v}_1^T (d_{...j})$, $d_{\mathbf{v}_{2j}} = \mathbf{v}_2^T (d_{...j})$, aso.
- interpretation:
 - the “new” x-coordinate of some data item
is a linear combination of all original attributes of this item,
weighted by the entries of eigenvector #1,
i.e.,
a large absolute component # k in eigenvector # g
tells about a major influence of data attribute # k
onto the principal component # g

PCA – Investigate the Loadings (3)



Back to our example:

- looking (again) at \mathbf{V} , we see:

$$\mathbf{V} = \begin{pmatrix} -0.7071 & 0.7071 \\ -0.7071 & -0.7071 \end{pmatrix}$$

- the 1st principal component is a (negative) combination of Height and Shoesize, whereas
- the 2nd principal component expresses their disagreement (e.g., large feet, but short)
- taking also the vastly different eigenvalues (~ 1.93 vs. ~ 0.07) into account, we find that
 - mostly, it's about the overall size (height & feet)
 - minor differences exist (larger feet, while being a bit shorter)

PCA – all of the data (1)



Considering all of the people's data, i.e.,

- the four data attributes Age, Height, Weight, and Shoesize

Steps: [\[see the MATLAB example 2d_F\]](#)

1. load the data
 2. set up the data **D** (here data cols., again, i.e., items in the rows);
compute also n , d , all per-dim. means **mu** & std.-dev. values **sigma**
 3. normalize the data: **Dz** becomes the z-score of **D**
 4. visualize to check
- [...]

PCA – all of the data (1)

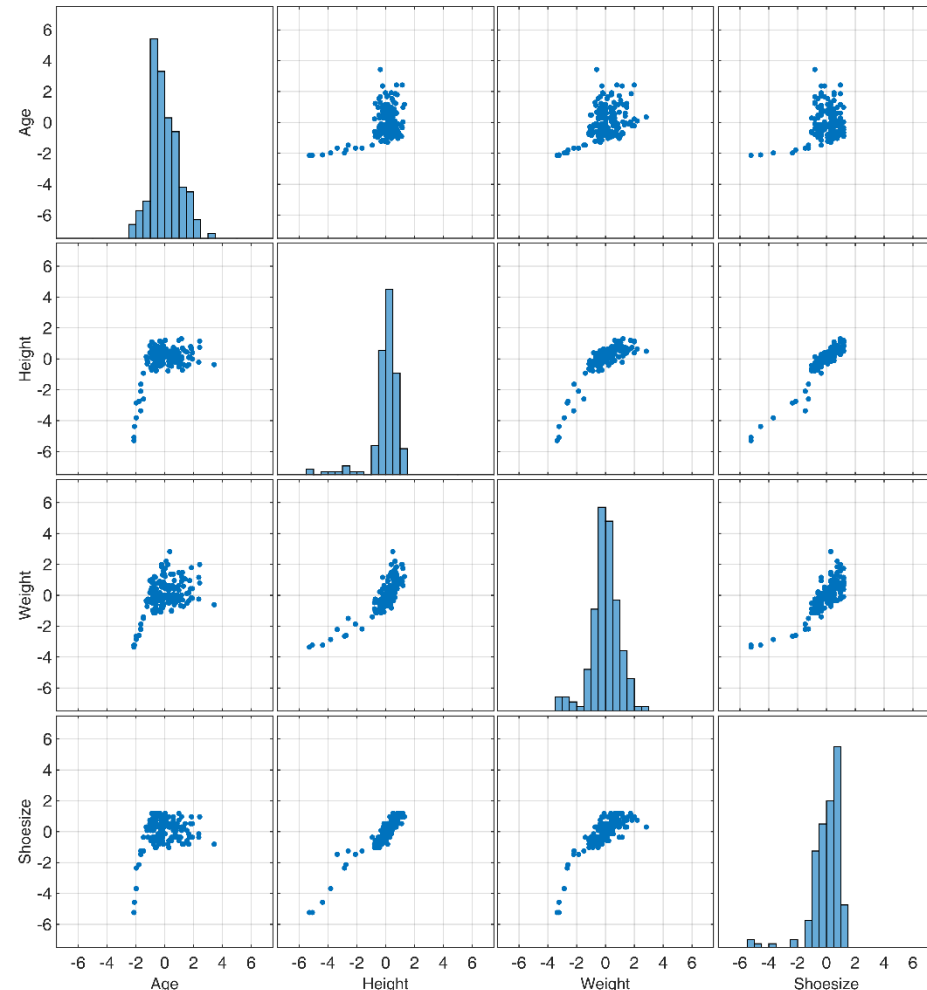


Considering all of the people's data, i.e.,

– the four data attributes Age, Height, Weight, and Shoesize

Steps:

1. load the data
2. set up the data **D** (here da compute also n , d , all per-
3. normalize the data: **Dz** b
4. visualize to check
- [...]



[see the MATLAB example 2d_F]

PCA – all of the data (2)



Considering all of the people's data, i.e.,

- the four data attributes Age, Height, Weight, and Shoesize

Steps:

[...]

5. compute the covariance matrix **S_{2z}**

and do the eigendecomposition into **V_z** and **L_z**

6. sort by the absolute eigenvalue value to get **$L_z S$** and **$V_z S$**

7. compute the scores **$D_z \text{Scores}$**

8. look at them

[...]

[see the [MATLAB example 2d_F](#)]

PCA – all of the data (2)



Considering all c

– the four data a

Steps:

[...]

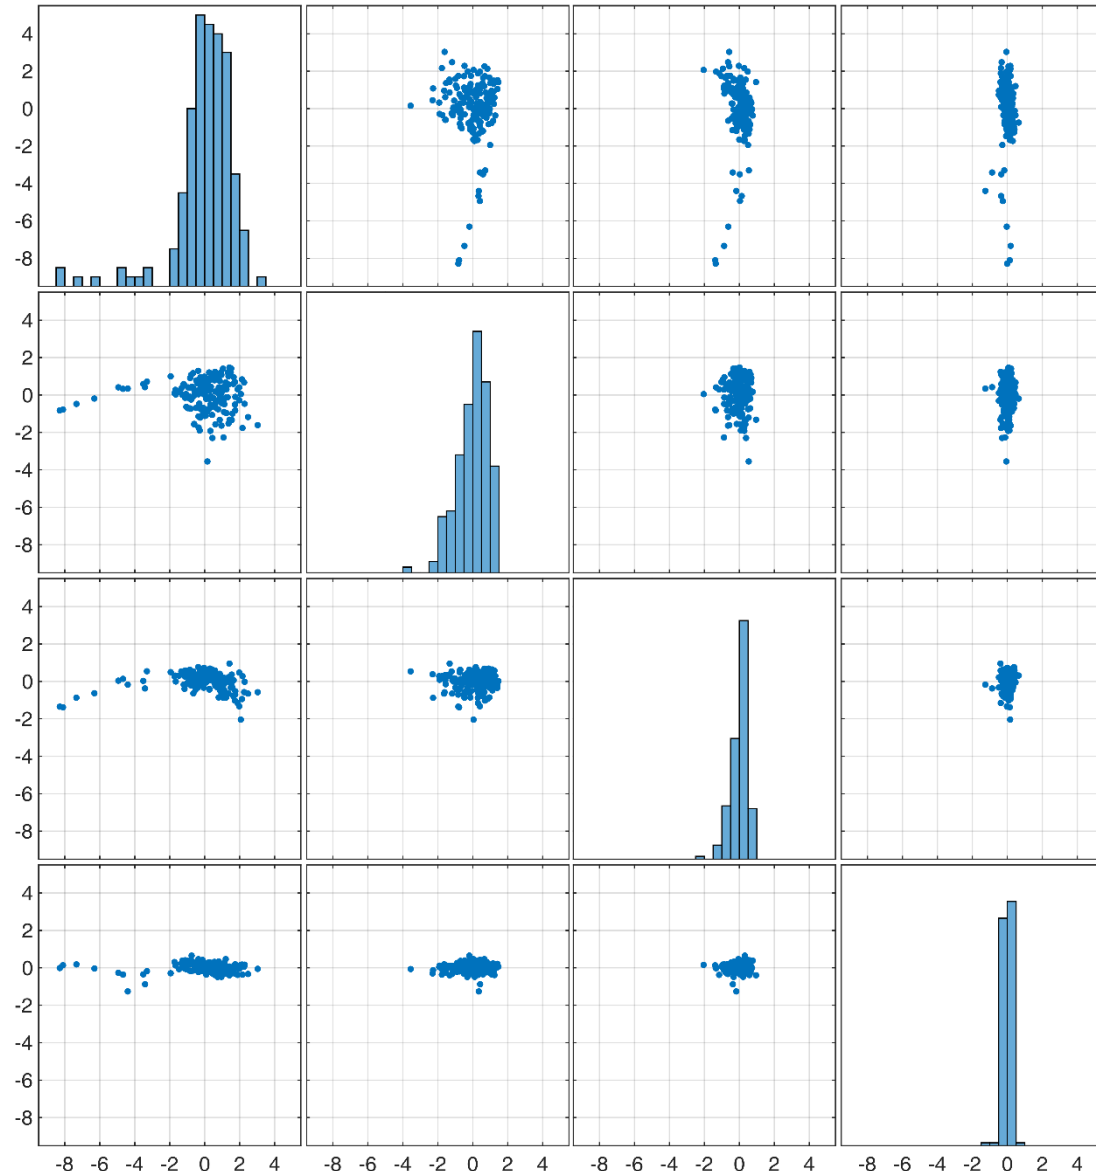
5. compute the cov
and do the eiger

6. sort by the abso

7. compute the scc

8. look at them

[...]



[see the MATLAB example 2d_F]

PCA – all of the data (3)



Considering all of the people's data, i.e.,

- the four data attributes Age, Height, Weight, and Shoesize

Steps:

[...]

9. do a 2D-reconstruction **DzReco2**

10. do a comparative visualization

11. undo the normalization for the reconstruction, get **Dreco2**

12. visualize

[...]

[see the MATLAB example 2d_F]

PCA – all of the data (3)

Considering all

– the four data

Steps:

[...]

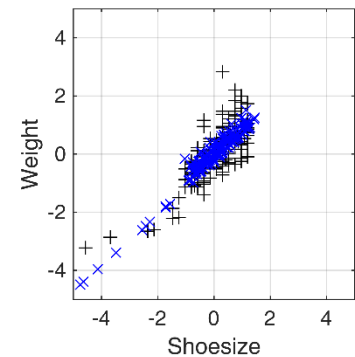




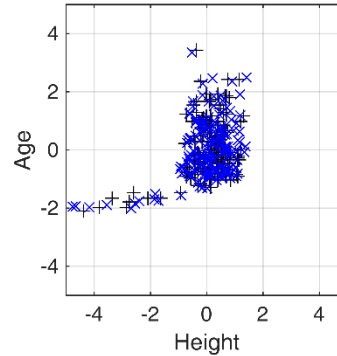
9. do a 2D-reco

10. do a compar

11. undo the nor

12. visualize

[...]



[see the MATLAB example 2d_F]

PCA – all of the data (3)



Considering all of the people's data, i.e.,

- the four data attributes Age, Height, Weight, and Shoesize

Steps:

[...]

9. do a 2D-reconstruction **DzReco2**

10. do a comparative visualization

11. undo the normalization for the reconstruction, get **Dreco2**

12. visualize

[...]

[see the [MATLAB example 2d_F](#)]

PCA – all of the data (3)



Considering all

– the four data

Steps:

[...]

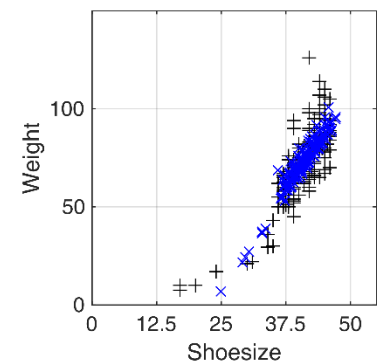





9. do a 2D-recor

10. do a compara

11. undo the norr

12. visualize

[...]



[see the [MATLAB example 2d_F](#)]

PCA – all of the data (4)



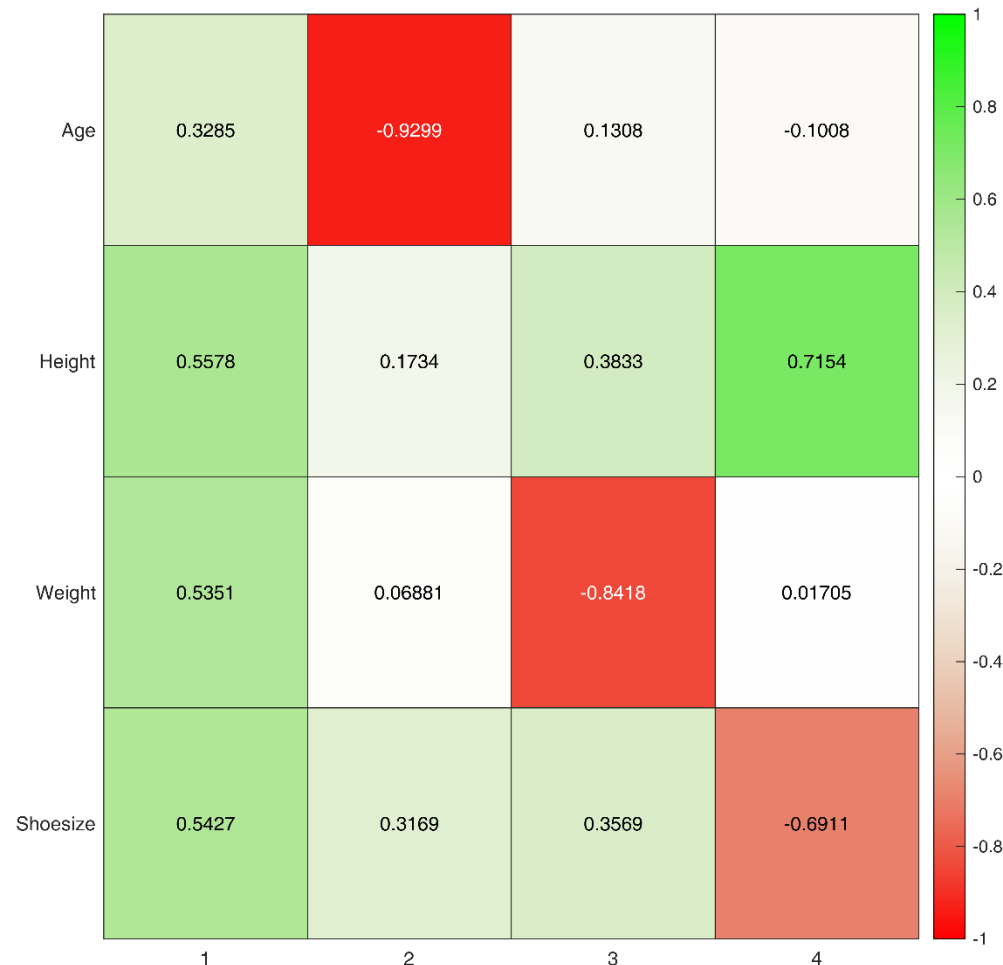
Considering all of the people's data, i.e.,

– the four data attributes Age, Height, Weight, and Shoesize

Steps:

[...]

13. look at the loadings

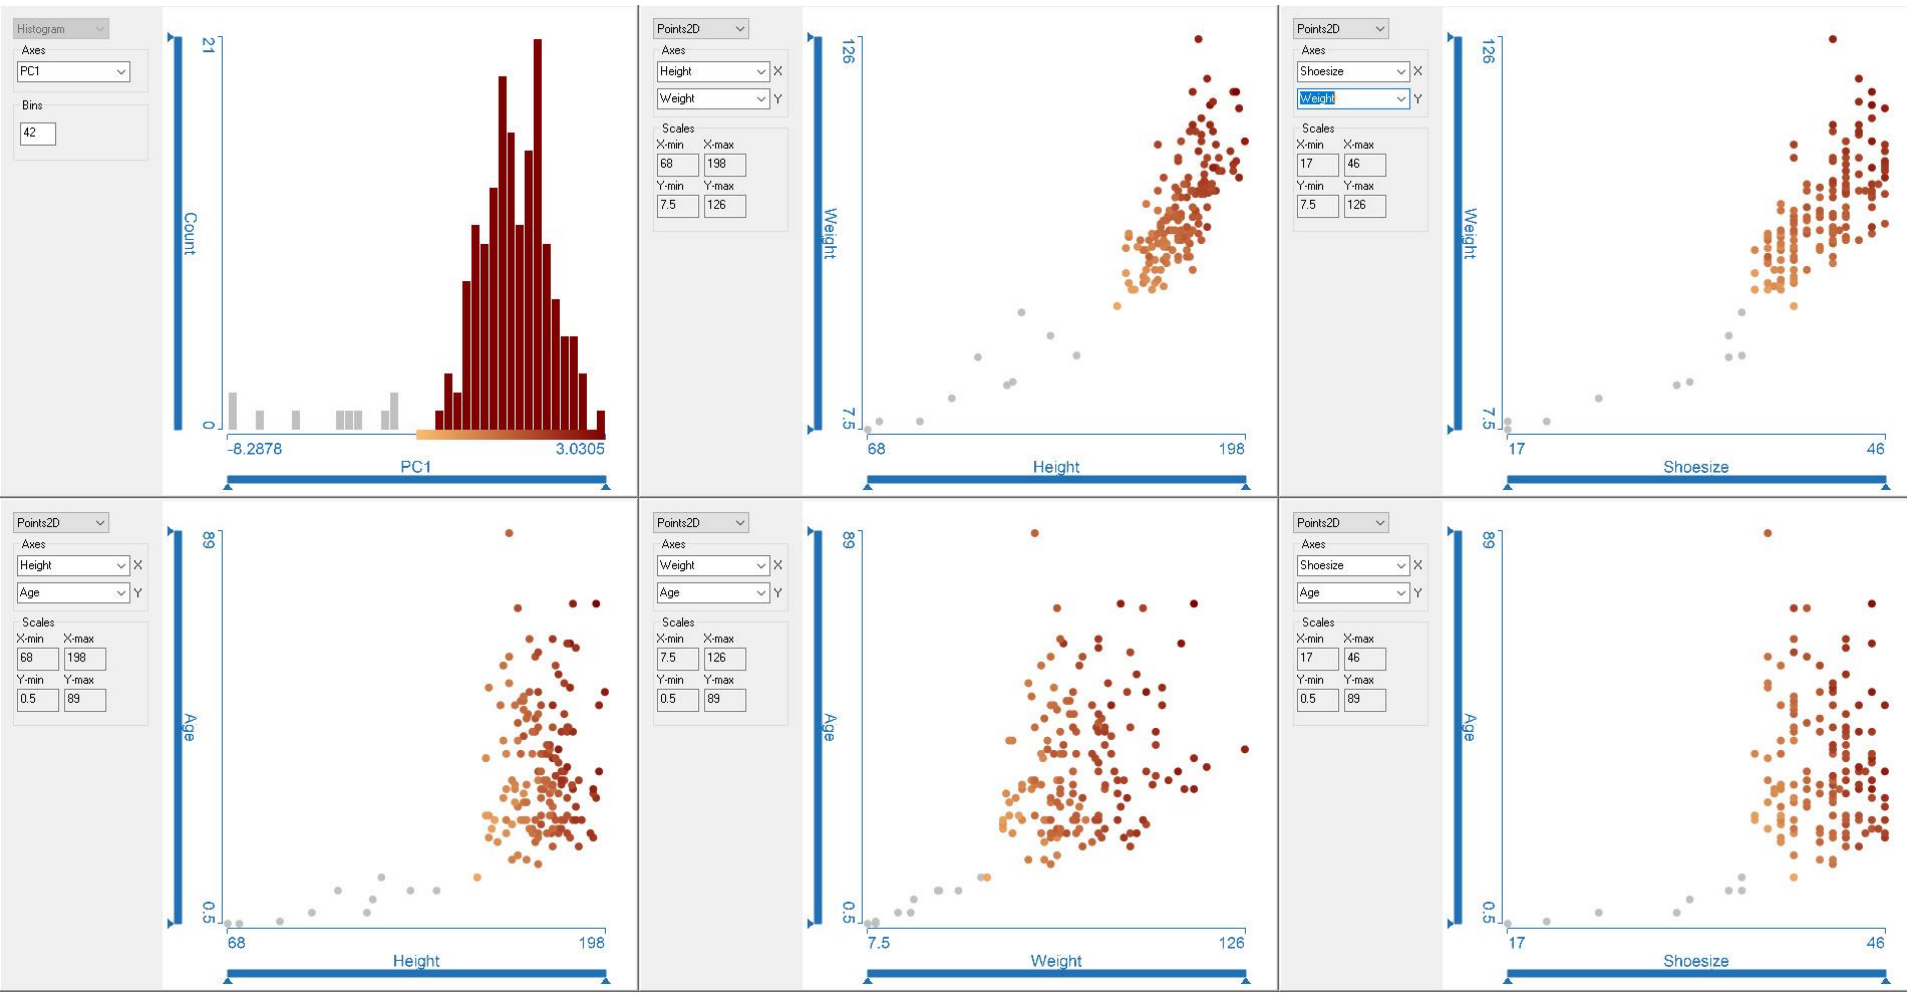
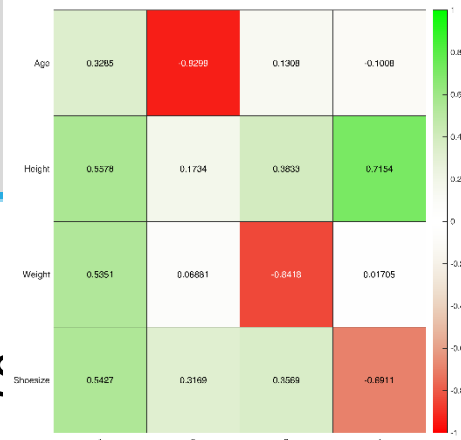


[see the MATLAB example 2d_F]

PCA – all of the data (4)

Considering all of the people's data, i.e.,








– the four data attributes Age, Height, Weight, and Shoesize



PCA – all of the data (4)

Considering all of the people's data, i.e.,

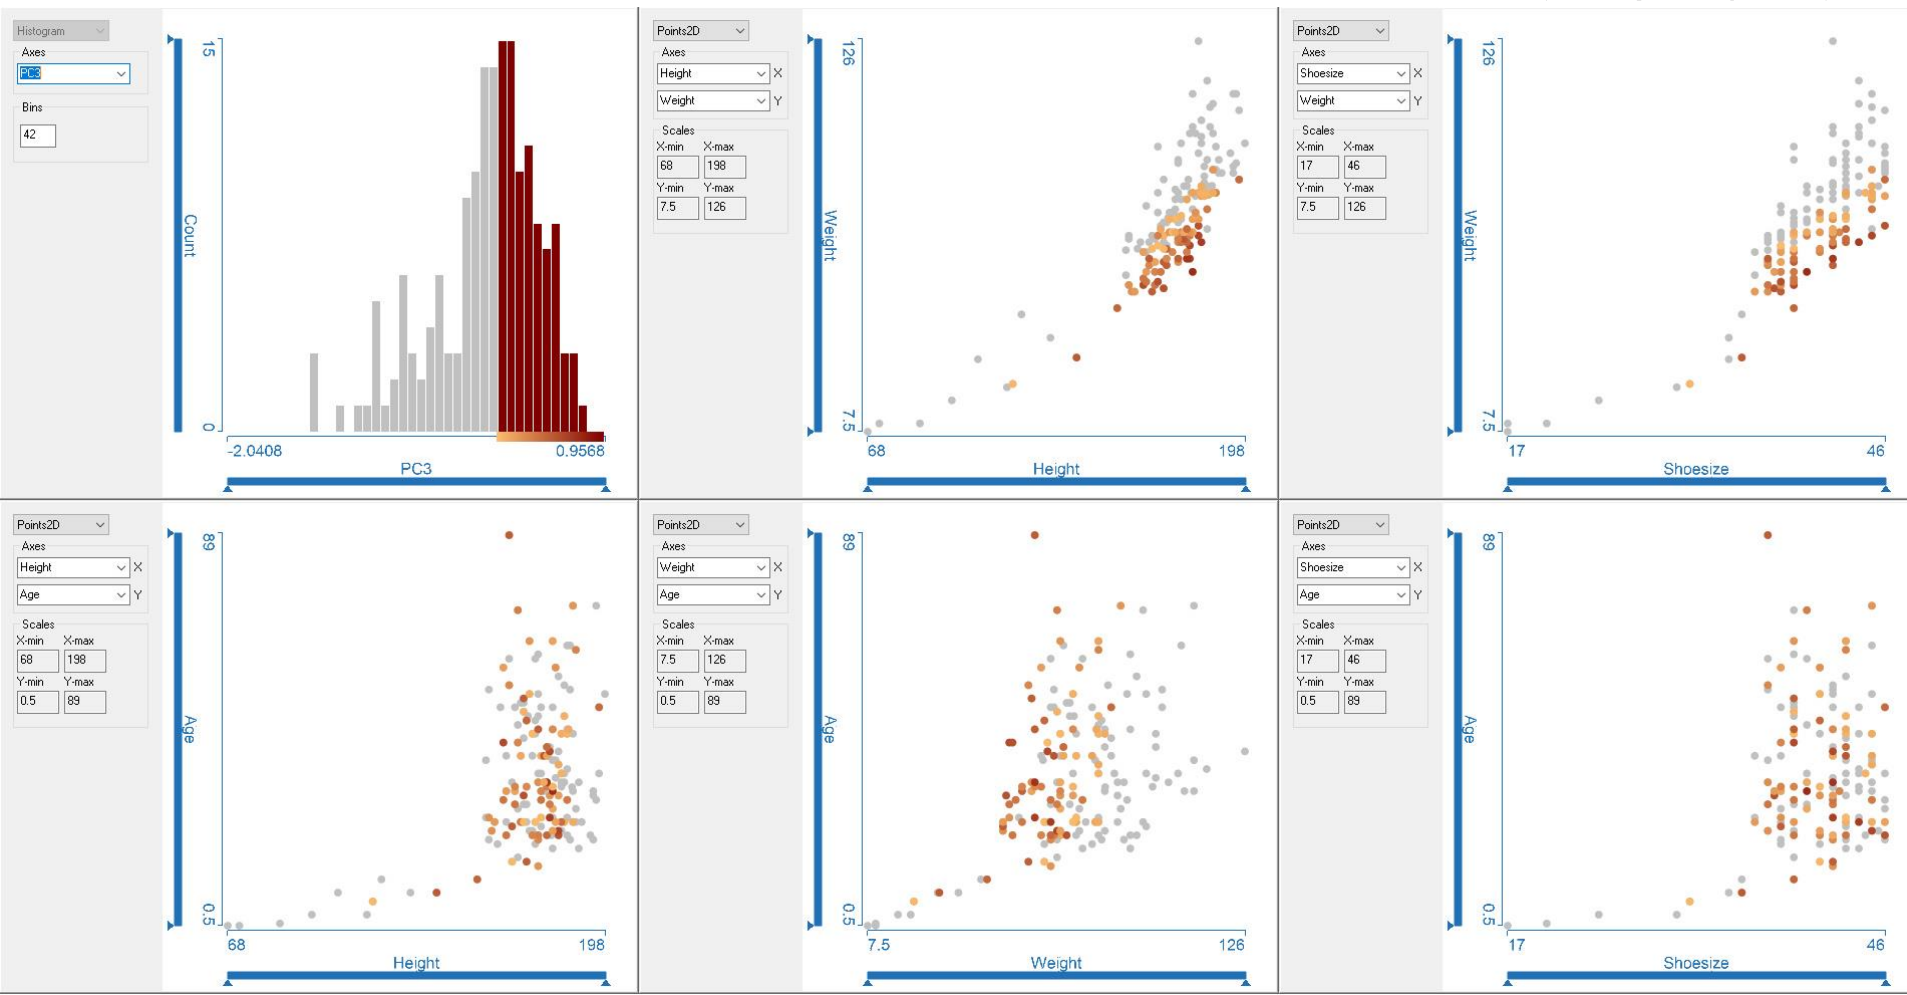

– the four data attributes Age, Height, Weight, and Shoesize



PCA – all of the data (4)

Considering all of the people's data, i.e.,

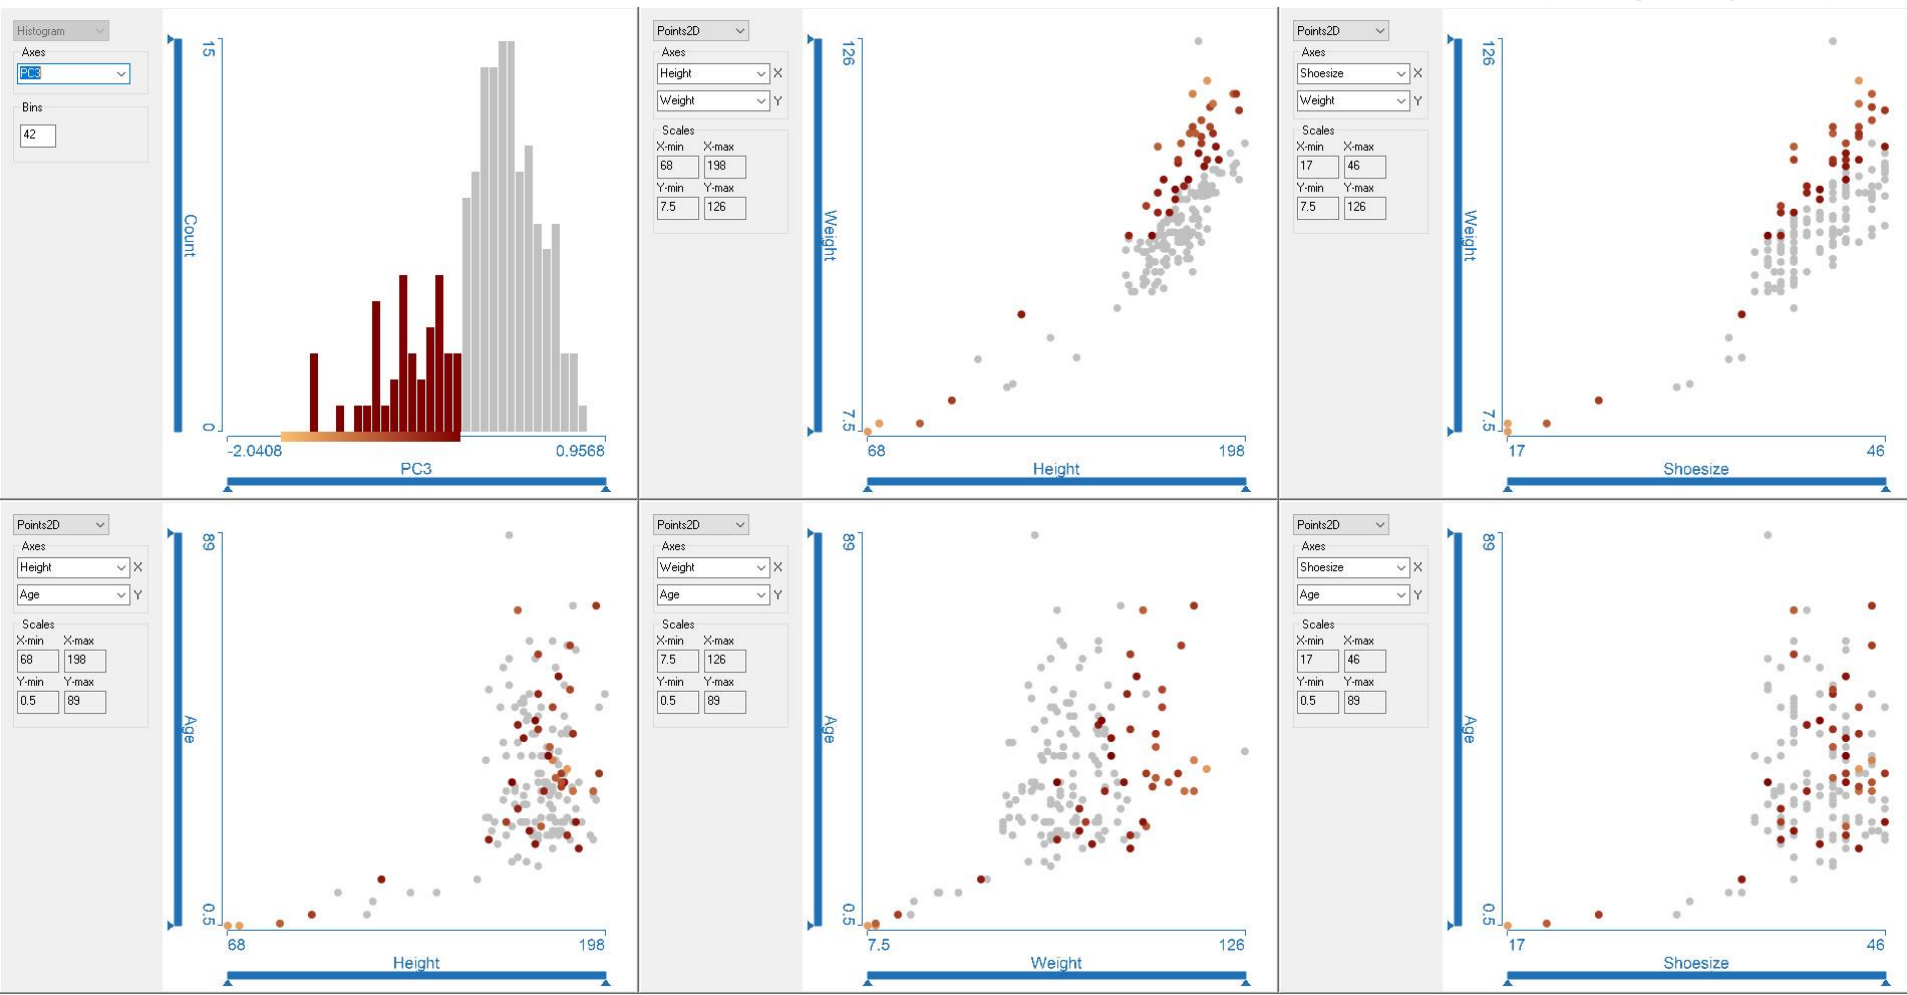

– the four data attributes Age, Height, Weight, and Shoesize



PCA – all of the data (4)

Considering all of the people's data, i.e.,

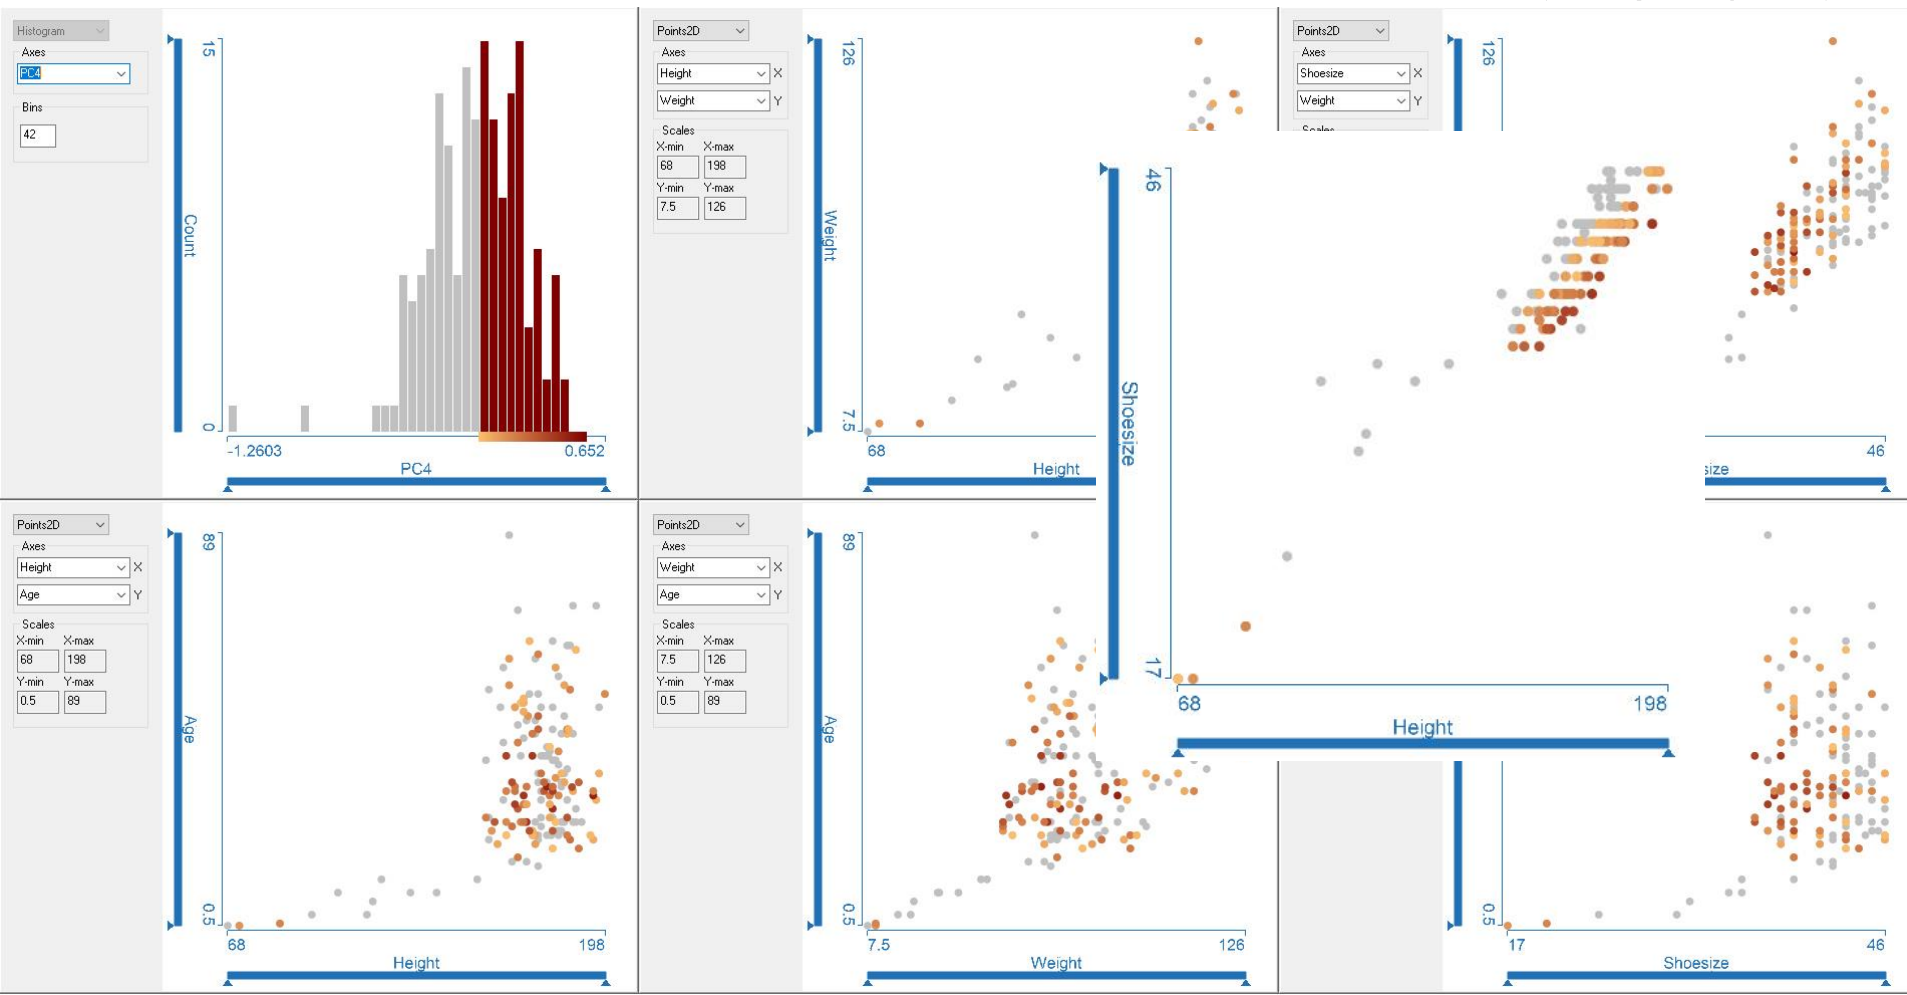
– the four data attributes Age, Height, Weight, and ShoeSize



PCA – all of the data (4)

Considering all of the people's data, i.e.,

– the four data attributes Age, Height, Weight, and Shoe size



Principal Component Analysis (PCA)

- can be seen as a ***best-possible linear technique for dimension reduction*** (best in the least sum of squared distances sense)
- is ***sensitive to the scaling of individual dimensions*** (therefore, often normalization is done dimension-wise first)
- this issue (PCA sensitive to dim.-scaling) implies that PCA is ***to a certain degree “arbitrary”, when dimensions of different units are studied*** (that cannot be mapped to each other, like temperatures and costs)
- PCA ***makes only sense, when the data is centered***
- numerous extensions / alternative approaches exist, including non-linear techniques, sparse PCA, robust PCA, ...

How to Solve an Eigenproblem

Focus up to here:

- how to use an eigenanalysis

Not touched (yet):

- how to solve an eigenproblem

In short:

- given $\mathbf{A} \mathbf{x} = \lambda \mathbf{x}$,
- we get to $(\mathbf{A} - \lambda \mathbf{I}) \mathbf{x} = 0$
 - if $(\mathbf{A} - \lambda \mathbf{I})$ is non-singular, then only $\mathbf{x} = \mathbf{0}$ is a solution
 - thus, and in order to find non-degenerated solutions, we need to require that $(\mathbf{A} - \lambda \mathbf{I})$ is singular, i.e., $\det(\mathbf{A} - \lambda \mathbf{I}) = 0$
- $\det(\mathbf{A} - \lambda \mathbf{I}) = 0$ is called the characteristic equation of \mathbf{A} , leading to the eigenvalues λ_i
- having the λ_i , we can solve $(\mathbf{A} - \lambda_i \mathbf{I}) \mathbf{x} = 0$ for the \mathbf{x}_i

Next:

- a new part of INF250 with
 - data fitting (somehow a continuation of this part!)
 - splines
 - iterative methods

Reading and Related Material

In the book:

- chapter 2 (on linear systems), mostly section 2.1, and related parts
- chapter 15 (on SVD, etc.)

Course notes:

- section 4

An interactive online-book:

- <http://ImmersiveMath.com/>,
chapters (1–4 &) 5 (on Gaussian elimination)
and 6 (on The Matrix)

On Wikipedia:

- many good pages